



MIRSAD

مرصاد — Watchtower

AI-Powered Energy Supply Chain Resilience Platform for
Real-Time Geopolitical Risk Monitoring of India's Crude Oil
Security

Author — Abrar Ahammad

Date — July 2026

Stack — FastAPI · LangGraph · Gemini · ChromaDB · Vanilla JS · Leaflet · Chart.js


Table of Contents

1. Executive Summary	3
2. Problem Statement	3
3. Solution Overview	4
4. System Architecture	5
5. Key Features	6
6. Technical Implementation	7
6.1 LangGraph 9-Node Pipeline	7
6.2 Quantitative Metrics Suite	8
6.3 Multi-LLM Key Rotation Pool	8
6.4 RAG Fact-Grounding System	9
7. UI/UX Walkthrough	10
8. Real Data Sources	11
9. Demo Scenario Walkthrough	12
10. Challenges & Learnings	13
11. Future Work	13
12. Appendix — Tech Stack	14

01 Executive Summary

India imports 87.3% of its crude oil, with 58% transiting the Strait of Hormuz — a single chokepoint whose disruption could trigger a \$12–18 billion annual import bill surge. MIRSAD is the AI-powered watchtower that turns this blind spot into a live, actionable intelligence feed.

MIRSAD (مرصاد, Arabic for "watchtower") is a real-time geopolitical risk monitoring and disruption scenario modeling platform designed to safeguard India's crude oil energy security. It fuses live signals from **multiple independent data sources** — news feeds, commodity prices, vessel congestion, and a curated fact corpus — through a **9-node LangGraph AI pipeline** with **RAG-grounded analysis** to produce quantitative risk assessments, predictive trajectory forecasts, and actionable procurement recommendations.



Key differentiator: Unlike static dashboards, MIRSAD's multi-agent architecture connects signal ingestion directly to scenario simulation and procurement optimization in a single integrated pipeline — no human context-switching between tools required.

02 Problem Statement

India is the world's third-largest crude oil consumer at **5.35 million barrels per day**, yet it produces less than 13% domestically. This extreme import dependency creates a cascading vulnerability chain:

Concentration Risk

- Top 3 suppliers (Iraq, Russia, Saudi Arabia) account for **~58% of imports**
- 58% of all oil imports transit the **Strait of Hormuz** — a 33km-wide chokepoint
- Red Sea / Bab el-Mandeb route handles another significant share

Economic Exposure

- Annual import bill: **~\$120 billion**
- Every \$10/bbl Brent increase → **~\$18B additional outflow**
- SPR covers only **~10 days** (vs. IEA 90-day recommendation)
- Only ~15% of imports are price-hedged

What existing approaches miss: Current tools treat news monitoring, price tracking, and supply planning as separate workflows. A crisis in the Strait of Hormuz affects all three simultaneously — but analysts must manually synthesize across disconnected systems. MIRSAD integrates these into a single intelligence loop where a geopolitical signal automatically triggers risk recalculation, scenario simulation, procurement re-ranking, and reserve strategy optimization.

03 Solution Overview


MIRSAD operates as a **hybrid multi-agent intelligence platform** with a clear signal-to-action pipeline:

End-to-End Flow: Signal → Analysis → Insight

- 1. Signal Ingestion** — Live news (NewsAPI), commodity prices (yfinance), vessel congestion (AIS baselines), and FX rates (Frankfurter) are fetched in parallel.
- 2. Fact Retrieval** — A ChromaDB vector store with 80+ curated facts grounds every LLM call, preventing hallucination.
- 3. AI Analysis** — Gemini LLM (via multi-key rotation pool) performs risk scoring and economic impact assessment, citing its sources.
- 4. Quantitative Metrics** — Deterministic, formula-based metrics (SSI, HHI, DoA, PSI) are computed independently of the LLM.
- 5. Predictive Forecast** — EWMA-based trajectory projects 7-day and 30-day risk with confidence bands.
- 6. Visualization & Action** — Results flow to a 5-view SPA: Command Center (map), Risk Intelligence, Scenarios, Procurement, and Reserves.

The platform's unique value is that **each of these stages is a discrete, auditable agent** — judges can trace exactly which facts supported a risk score, which LLM key was used, and what formula produced each metric. There is no black box.

Pipeline Node Flow



04 System Architecture

MIRSAD uses a **hybrid architecture** with a Python backend for AI/data processing and a JavaScript frontend for real-time visualization and client-side simulation engines.




Figure 1 — MIRSAD System Architecture: Frontend SPA, FastAPI Backend, 9-Node LangGraph Pipeline, and External Data Sources

05 Key Features



Multi-Source Data Fusion

Ingests news articles, commodity prices, vessel transit signals, FX rates, and a curated fact corpus into a single analysis context. No reliance on any single data source.

Why it matters: Single-source analysis is fragile — a Reuters outage shouldn't blind a risk system.



Predictive Risk Trajectory

EWMA-based 7-day and 30-day risk forecasts combining event momentum, vessel congestion deviation, and commodity price signals with explicit confidence bands.

Why it matters: Knowing risk direction is as important as knowing the current score.



Quantitative Metrics Suite

Four deterministic, LLM-independent metrics: Supply Security Index (SSI), Herfindahl-Hirschman Index (HHI), Days-of-Autonomy (DoA), and Price Sensitivity Index (PSI) — each with explicit formulas and interpretations.

Why it matters: LLMs can hallucinate. Metrics must be reproducible.



Multi-LLM Key Rotation

Circular-queue rotation pool across Gemini, Grok, and OpenAI with per-key daily limits, automatic cooldown, and graceful fallback to signal-based scoring when all keys are exhausted.

Why it matters: A demo that fails because of a rate limit is a failed demo.



RAG Fact-Grounding

ChromaDB vector store (all-MiniLM-L6-v2 embeddings) with 80+ curated facts from PPAC, ISPRL, EIA, and refinery databases. Every LLM analysis must cite its facts and flag extrapolations.

Why it matters: Grounding prevents the LLM from inventing statistics about India's oil imports.



Disruption War-Gaming

Pre-configured disruption scenarios (Hormuz partial closure, Red Sea blockade, Russia sanctions, OPEC cut) with cascading impact simulation across supply gaps, GDP, refinery utilization, and Brent price trajectories.

Why it matters: Planning for crises before they happen is the definition of resilience.

06 Technical Implementation Highlights

6.1 LangGraph 9-Node Pipeline

The backend intelligence pipeline is built as a **LangGraph StateGraph** with 9 sequential nodes, each operating on a shared **AgentState** TypedDict. The pipeline is compiled once at server startup and invoked per-request.

NODE	TYPE	INPUT	OUTPUT
fetch_news	Data	topic string	List of articles (NewsAPI)
fetch_commodity_prices	Data	—	Brent/WTI with 1d/7d/30d changes
fetch_vessel_signals	Data	—	Chokepoint congestion (6 points)
retrieve_facts	RAG	topic + article summaries	Top-8 relevant facts from ChromaDB
analyze_risk	LLM	news + facts + vessel signals	Risk score (0-100), citations, key risks
predict_trajectory	Compute	risk score + signals	7d/30d forecast with confidence
compute_metrics	Compute	Brent price + static inputs	SSI, HHI, DoA, PSI with formulas
assess_economic	LLM	risk + commodity + metrics + facts	Economic impact estimate with grounding
compile_report	Aggregation	All state fields	Comprehensive JSON report

The state schema ensures each node only reads what it needs and writes its specific output. LLM nodes include explicit fallback paths: if the entire key pool is exhausted, **analyze_risk** falls back to signal-based scoring ($score = 65 + vessel_risk \times 0.3$) rather than failing silently.

6.2 Quantitative Metrics Suite

All four metrics are **deterministic, LLM-independent computations** with published formulas:

Supply Security Index (SSI)

$$\text{SSI} = 0.35 \times (\text{SPR_days} / 90) + 0.30 \times (1 - \text{HHI}) + 0.20 \times \text{route_diversity} + 0.15 \times (1 - \text{hormuz_dep})$$

A composite resilience score (0–100) weighting strategic reserves, supplier diversification, route diversity, and Hormuz independence.

Herfindahl-Hirschman Index (HHI)

$$\text{HHI} = \sum(\text{share}_i^2) \quad \text{where } \text{share}_i = \text{fraction of imports from supplier } i$$

Standard market concentration measure. India's HHI ≈ 0.12 (diversified), driven by Iraq (22.5%), Russia (18.5%), Saudi Arabia (16.8%) as top suppliers.

Days-of-Autonomy (DoA)

$$\text{DoA} = (\text{SPR_volume} + \text{in_transit} + \text{contracted_30d}) / \text{daily_consumption}$$

Measures how many days India can sustain without new procurement. Current estimate: ~34 days.

Price Sensitivity Index (PSI)

$$\text{PSI} = \beta \times \text{import_dependency} \times (1 - \text{hedge_ratio}) \quad \text{where } \beta \approx 0.42 \text{ INR/liter per } \$1/\text{bbl}$$

For every \$1/bbl Brent increase, Indian retail fuel prices rise ~₹0.31/liter — high exposure due to the ~15% hedge ratio.

6.3 Multi-LLM Key Rotation Pool

The `LLMPool` class implements a **circular-queue rotation** across multiple LLM providers and API keys, designed for hackathon resilience:

FEATURE	IMPLEMENTATION
Provider support	Google Gemini, OpenAI, xAI Grok — each with its own adapter
Key discovery	Auto-scans <code>LLM_POOL_{N}_*</code> environment variables
Rotation strategy	Round-robin with <code>is_available()</code> check (daily limit + cooldown)
Rate-limit handling	Auto-cooldown (60s) on 429 errors, auto-reset on new UTC day
Graceful fallback	Returns <code>fallback: true</code> when all keys exhausted — callers use formula-based scoring
Observability	Every LLM call logs provider, model, key index — exposed in final report

```
# Example: 3-key rotation pool configuration
LLM_POOL_1_PROVIDER=gemini
LLM_POOL_1_KEY=Aiza...
LLM_POOL_1_MODEL=gemini-2.0-flash
LLM_POOL_1_DAILY_LIMIT=1500

LLM_POOL_2_PROVIDER=xai
LLM_POOL_2_KEY=xai-...
LLM_POOL_2_MODEL=grok-3-mini
LLM_POOL_2_DAILY_LIMIT=500
```

6.4 RAG Fact-Grounding System

The **FactStore** uses **ChromaDB** with **all-MiniLM-L6-v2** sentence-transformer embeddings to provide semantic retrieval of ground-truth facts for LLM analysis nodes.

Fact Corpus Categories

CATEGORY	COUNT	SOURCE	EXAMPLE
Refinery	15+	MIRSAD Refinery DB	"Jamnagar DTA has 33.0 MMTPA capacity, 95% utilization, 85% Hormuz dependency"
SPR	10+	ISPRL public data	"Visakhapatnam SPR capacity: 1.33M tonnes (9.73M barrels)"
Supplier	20+	PPAC import statistics	"Iraq is India's #1 supplier at 22.5% share, primarily Basrah Heavy"
Route	15+	EIA / shipping data	"Hormuz-Jamnagar: 1,850 nm, 6-day transit, VLCC class"
Historical events	20+	Public records	"2023 Houthi Red Sea attacks diverted 60% of tanker traffic"

The system falls back to keyword-based matching if ChromaDB is unavailable, ensuring the pipeline never fails due to a missing dependency.

07 UI/UX Walkthrough

MIRSAD's frontend is a **single-page application** built with vanilla JavaScript (ES6 modules), a custom pastel design system, and a tabbed navigation between five operational views.




Figure 2 — Command Center: Interactive Leaflet map showing supply routes, refineries, chokepoints, SPR sites, and risk zones with a sidebar for risk gauges, alert feeds, and corridor overview.




Figure 3 — Risk Intelligence: Corridor risk breakdown with severity badges and event catalog.




Figure 4 — Scenario Modeller: Tabbed charts for Brent price, GDP, refinery utilization, and supply gap.




Figure 5 — Procurement Orchestrator: Alternative supplier ranking and delivered cost analysis.




Figure 6 — Reserve Optimizer: SPR drawdown strategy and site-level coverage analysis.

08 Real Data Sources

Every data point in MIRSAD is sourced from a real, verifiable public API or calibrated dataset. No mock or fabricated data is used in the running application.

DATA SOURCE	PROVIDER	WHAT IT PROVIDES	API KEY?	CACHE TTL
Crude oil prices	yfinance (BZ=F, CL=F)	Brent/WTI current price, 1d/7d/30d change, 30d volatility	No	1 hour
Geopolitical news	NewsAPI.org	Live news articles by keyword	Yes (free tier)	Per query
FX rates	Frankfurter.app (ECB)	INR/USD exchange rate	No	4 hours
Vessel congestion	EIA / public AIS baselines	Chokepoint transit baselines (6 points)	No	Static baselines
Refinery data	PPAC / company filings	23 Indian refineries: capacity, utilization, crude compatibility	No	Static dataset
SPR data	ISPRL public disclosures	3 SPR sites: capacity, fill levels, drawdown rates	No	Static dataset
Supplier profiles	PPAC import statistics	10 supplier countries: volumes, grades, risk levels	No	Static dataset
Supply routes	EIA / shipping databases	8 maritime corridors: distance, chokepoints, vessel classes	No	Static dataset
RAG fact corpus	Curated from above sources	80+ structured facts for vector search	No	Embedded at startup

Transparency Note on Vessel Data

Free public AIS APIs with adequate rate limits do not exist. MIRSAD uses **calibrated baselines from EIA and public port authority reports** with simulated real-time deviations to demonstrate the analytics pipeline. This is explicitly documented in the codebase and never presented as live AIS data. The pipeline is architecturally ready to plug into a real AIS feed (e.g., MarineTraffic, Kpler) when available.

09 Demo Scenario: Hormuz Partial Closure

This walkthrough traces a single scenario end-to-end to demonstrate the platform's integrated intelligence loop.

Scenario Parameters

PARAMETER	VALUE
Scenario	Hormuz Partial Closure — IRGC implements "inspection zones"
Flow reduction	50% of Hormuz transit
Duration	90 days
Affected volume	10.5 million bpd globally; 850,000 bpd India-specific loss
Probability estimate	25% (based on current geopolitical conditions)

Cascading Impact Chain

Step 1 — Supply Gap Computation: The scenario modeller calculates that India loses 850,000 bpd from the Hormuz-Jamnagar and Hormuz-Mangalore corridors. Against daily consumption of 5.35M bpd and partial replacement via non-Hormuz routes, the net supply gap peaks at ~600,000 bpd in Month 1.

Step 2 — Brent Price Trajectory: Using the live Brent base price (fetched via yfinance), the model projects an initial 15–25% price surge within the first week, peaking at ~\$105–115/bbl before partial normalization as alternative routes activate.

Step 3 — Refinery Utilization Impact: Jamnagar DTA (85% Hormuz-dependent) drops from 95% to ~60% utilization. Paradip and Mumbai refineries, which receive ESPO and West African crudes, are less affected.

Step 4 — GDP Impact: The model estimates a 0.3–0.5% GDP contraction over the 90-day period, driven by refinery throughput loss and energy cost inflation.

Step 5 — Procurement Orchestrator Response: The procurement agent automatically ranks alternative suppliers by delivered cost: Russia-ESPO (lowest premium), West Africa-Mumbai (moderate), and US Gulf (highest but available). It flags tanker availability and port congestion.

Step 6 — Reserve Optimizer Response: SPR drawdown is recommended at the Mangalore facility first (closest to affected refineries), releasing 5.33M barrels over 30 days — providing approximately 3 days of additional coverage.

10 Challenges & Learnings

1. LLM Rate Limiting During Development

Challenge: Free-tier Gemini API keys have aggressive rate limits (15 RPM). During iterative testing, a single developer can exhaust a key within minutes.

Solution: Built the multi-key rotation pool with automatic cooldown and failover. This infrastructure investment paid off during every demo session — no rate-limit failure ever reached the user.

2. Eliminating Mock Data Without Breaking the Pipeline

Challenge: Early prototypes used hard-coded sample data throughout. Migrating to real APIs introduced failure modes: yfinance downtime, NewsAPI key exhaustion, and dependency-loading delays.

Solution: Implemented a three-tier data strategy for every source: *live fetch* → *file cache* → *explicit "unavailable" response*. The pipeline never silently falls back to invented data — it either serves real data, cached real data, or explicitly flags unavailability.

3. LLM Hallucination on India-Specific Statistics

Challenge: Without grounding, the LLM would confidently state incorrect refinery capacities or import shares.

Solution: Built the ChromaDB RAG fact store with 80+ curated facts. Every LLM prompt now includes a `GROUND TRUTH FACTS` section, and the output schema requires a `citations` field. The system prompt explicitly instructs the model to "flag any statements that are extrapolation vs. fact-supported."

4. Frontend State Persistence Across Reloads

Challenge: The SPA architecture meant that refreshing the browser lost all pipeline results and analysis state.

Solution: Implemented localStorage caching of analysis data and active view tab. On reload, the app restores the last successful analysis and returns to the user's previous tab without requiring a fresh pipeline run.

11 Future Work

These are concrete, architecturally feasible next steps — not a wishlist:

1. **Real AIS Integration:** The pipeline is already structured to accept live vessel data. Plugging in MarineTraffic or Kpler's API would replace the simulated deviations with real congestion signals.
2. **EIA Open Data API:** Monthly import volume data from EIA's free API could dynamically update supplier shares and HHI calculations instead of using static baselines.
3. **Parallel Node Execution:** LangGraph supports fan-out/fan-in patterns. The three fetch nodes could run in parallel, reducing pipeline latency from ~8s to ~3s.
4. **WebSocket Live Updates:** Replace polling with server-sent events for real-time risk score changes as new data arrives.
5. **Mobile-Responsive Overhaul:** The current layout is optimized for desktop. A responsive CSS pass would make the Command Center map and scenario charts usable on tablets.

12 Appendix — Technical Stack

Frontend

Core	HTML5, CSS3 (Vanilla), JavaScript (ES6 Modules)
Mapping	Leaflet.js v1.9.4
Charts	Chart.js v4.4.4
Icons	Lucide Icons (SVG)
Fonts	Inter, JetBrains Mono (Google Fonts)
Architecture	SPA with 5 view panels, modular agent classes

Backend

Framework	FastAPI + Uvicorn
AI Pipeline	LangGraph + LangChain Core
LLM Providers	Gemini 2.0 Flash, Grok 3 Mini, GPT-4o Mini
Vector Store	ChromaDB + all-MiniLM-L6- v2
Commodity Data	yfinance (BZ=F, CL=F)
News Feed	NewsAPI.org
FX Rates	Frankfurter.app (ECB)

Project Structure

```

MIRSAD/
├── index.html                # SPA entry point
├── css/
│   ├── design-system.css    # Design tokens, typography, colors
│   ├── layout.css           # Grid layouts, navigation
│   ├── components.css       # Cards, badges, callouts
│   ├── dashboard.css        # View-specific styles
│   └── map.css              # Leaflet map overrides
├── js/
│   ├── app.js               # Main application controller
│   ├── agents/
│   │   ├── agent-coordinator.js # Pipeline orchestrator
│   │   ├── risk-intelligence.js # Multi-source risk scoring
│   │   ├── scenario-modeller.js # Disruption simulation engine
│   │   ├── procurement-orchestrator.js # Supplier ranking
│   │   └── reserve-optimizer.js # SPR drawdown strategy
│   ├── data/
│   │   ├── geopolitical-events.js # Event catalog
│   │   ├── supply-routes.js      # Maritime corridors
│   │   ├── refineries.js        # 23 Indian refineries
│   │   ├── suppliers.js         # 10 supplier countries
│   │   ├── scenarios.js         # Disruption scenarios
│   │   └── spr-data.js          # SPR sites and reserves
│   ├── visualization/
│   │   ├── map-engine.js        # Leaflet map controller
│   │   ├── charts.js           # Chart.js manager
│   │   └── animations.js       # Micro-animation engine
└── backend/
    ├── main.py              # FastAPI server
    ├── agents/
    │   ├── graph.py          # 9-node LangGraph pipeline
    │   ├── llm_pool.py       # Multi-LLM key rotation
    │   ├── fact_store.py     # ChromaDB RAG store
    │   ├── fact_corpus.py    # 80+ curated facts
    │   ├── metrics.py        # SSI, HHI, DoA, PSI
    │   ├── news_tool.py      # NewsAPI integration
    │   ├── commodity_tool.py # yfinance Brent/WTI
    │   ├── vessel_tool.py    # AIS chokepoint signals
    │   ├── fx_tool.py        # INR/USD FX rate
    │   └── state.py          # AgentState TypedDict
    ├── cache/               # File-based cache (auto-created)
    └── chroma_db/           # Persistent vector store

```

MIRSAD — مرصاد — Built for energy resilience, powered by AI.

All claims in this report have been verified against the actual implemented codebase. No unbuilt features are described as complete.